

SHUBHODEEP MITRA

(602) 849-3716 | smitra27@asu.edu | [linkedin.com/in/shubhodeep-mitra](https://www.linkedin.com/in/shubhodeep-mitra) | [shubhodeepmitra.github.io](https://github.com/shubhodeepmitra)

EDUCATION

Master of Science in Computer Science

Arizona State University, Tempe, AZ

August 2024

GPA: 3.96/4

- Foundations of Algorithms, Distributed Database Systems, Privacy and Machine Learning, Cloud Computing, Statistical Machine Learning, Multimedia and Web Databases, and Blockchain Engineering.

Bachelor of Engineering in Computer Science and Engineering

The National Institute of Engineering, Mysuru, India

May 2018

GPA: 8.74/10

TECHNICAL SKILLS

Programming Languages: Java, C/C++, Bash, Python (NumPy, pandas), Android, Kotlin, Go, Rust.

Skills: SpringBoot, ReactJS, Redux, MapReduce, Kafka, Docker, REST, Ansible, AWS, GraphQL, PyTorch, Kubernetes, SQL, PostgreSQL, Redis, MongoDB, Firebase, Git, GitHub, MATLAB, A/B testing, Linux/Unix Programming, gdb, ElasticSearch.

EXPERIENCE

Graduate Research Assistant (Distributed Systems) - EMITLab ASU, Tempe, USA

December 2022 – Present

- Led development of scalable solutions for vertical scaling of large models, efficient data management, and concurrent workflows on AWS infrastructure, boosting sampling pool and simulation diversity by **4x**.
- Optimized deployment, configuration, and orchestration of a cluster of AWS EC2 using Ansible, driving a **27% reduction** in per-simulation experiment time through improved parallelism.
- Revamped in-memory graph data structure and integrated local caching mechanisms to generate provenance graph and multi-variate timelines, reducing MongoDB calls by over **50k**.

Software Engineer 2 - Hewlett Packard Enterprise (HPE), Bangalore, India

July 2018 – July 2022

- Led development efforts in C/C++ for critical modifications in internal-VLAN, L3 counters, Netdev, and Ofproto, pivotal for IP-Subinterface, facilitating seamless traffic flow across **17 protocols**.
- Redesigned IGMP and MLD multicast protocols to integrate real-time packet flow monitoring, enhancing reliability.
- Developed CLI infrastructure for L2 protocols, including VLAN, VLAN Translation, and Multi-Zone User-Based Tunneling, integrating SDN features to support advanced network management and configuration capabilities.
- Engineered a comprehensive test automation suite for all development work, achieving **98%** code coverage and empowering DevOps team to conduct periodic system health checks.
- Secured first place in a hackathon by addressing parking space challenges at HPE Bangalore campus.

RELEVANT PROJECTS

ColumnarDB

March 2024

- Designed a Columnar Database system in Java with features including BitMap and BTree indexing, compressed BitMap, Columnar Joins, Columnar Sort, Scan, and Delete optimizing data operations efficiency for **50k entries**.

Live Video Streaming Application

December 2023

- Architected and implemented a high-performance live video streaming application in **Go** and **WebRTC**, optimizing for low latency and scalability, employing advanced techniques to enhance video delivery quality.

Image Retrieval and Recommendation

October 2023

- Designed and built an end-to-end image retrieval and recommendation system with **82% accuracy**, applying neural network feature extraction, dimensionality reduction techniques, graph algorithms, and LSH.

Elastic Face Recognition Service

March 2023

- Developed face recognition services from videos leveraging **AWS IaaS** with 98% accuracy. Dynamically scaled in and out on-demand, efficiently handling concurrent user requests. Utilized AWS EC2, S3, DynamoDB, and SQS.

Scalable Aesthetic-Preserving Face De-Identification

November 2022

- Created a Kotlin Android app employing an **ML-kit** and openCV foundation to detect individuals in photograph backgrounds through face recognition, applying an aesthetic-preserving filter to safeguard bystanders' privacy.

Taxi Trajectory Visualization

October 2022

- Implemented an application enabling spatial-temporal, and KNN queries on NYC taxi data stored in Spark SQL table.

Real-Time Parking Spot Notification, Hewlett Packard Enterprise

January 2020

- Engineered a scalable system to monitor over 500 vehicles entering and exiting campus. Leveraged surveillance footage and employed OpenCV and TensorFlow, providing Android users with live parking availability updates.

INVOLVEMENT

Teaching Assistant, CSE510 - Database Management System Implementation

January 2024 – May 2024

Global Leadership Academy program, Arizona State University, Tempe, USA

August 2022 – December 2022

Novel TLS Signature Extraction for Malware Detection, IEEE CONECCT

July 2020

- Researched a solution at Hewlett Packard Enterprise to identify presence of malware in a network flow from an initial unencrypted Client Hello packet of TLS with 92.4 percent accuracy.